

ql8rmz3bn

December 14, 2023

Task -2

Perform data cleaning and exploratory data analysis (EDA) on a dataset of your choice, such as the Titanic dataset from Kaggle. Explore the relationships between variables and identify patterns and trends in the data.

Import necessary libraries

```
[ ]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[ ]: from google.colab import files
raw = files.upload()
```

<IPython.core.display.HTML object>

Saving titanic.csv to titanic.csv

```
[ ]: data = pd.read_csv("titanic.csv")
```

```
[ ]: df = data.copy()
```

```
[ ]: df
```

```
[ ]:
PassengerId  Survived  Pclass  \
0            1         0       3
1            2         1       1
2            3         1       3
3            4         1       1
4            5         0       3
..          ...      ...     ...
886          887         0       2
887          888         1       1
888          889         0       3
889          890         1       1
890          891         0       3
```

Name Sex Age SibSp \

0		Braund, Mr. Owen Harris	male	22.0	1
1		Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1
2		Heikkinen, Miss. Laina	female	26.0	0
3		Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1
4		Allen, Mr. William Henry	male	35.0	0
..	
886		Montvila, Rev. Juozas	male	27.0	0
887		Graham, Miss. Margaret Edith	female	19.0	0
888		Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1
889		Behr, Mr. Karl Howell	male	26.0	0
890		Dooley, Mr. Patrick	male	32.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S
..
886	0	211536	13.0000	NaN	S
887	0	112053	30.0000	B42	S
888	2	W./C. 6607	23.4500	NaN	S
889	0	111369	30.0000	C148	C
890	0	370376	7.7500	NaN	Q

[891 rows x 12 columns]

Exploratory data analysis

```
[ ]: df.shape
```

```
[ ]: (891, 12)
```

```
[ ]: df.head()
```

```
[ ]: PassengerId  Survived  Pclass  \
0                1         0       3
1                2         1       1
2                3         1       3
3                4         1       1
4                5         0       3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	

4	Allen, Mr. William Henry	male	35.0	0
---	--------------------------	------	------	---

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[ ]: df.tail()
```

```
[ ]:
      PassengerId  Survived  Pclass
886           887         0        2
887           888         1        1
888           889         0        3
889           890         1        1
890           891         0        3
      Name \
886      Montvila, Rev. Juozas
887      Graham, Miss. Margaret Edith
888      Johnston, Miss. Catherine Helen "Carrie"
889      Behr, Mr. Karl Howell
890      Dooley, Mr. Patrick
```

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
886	male	27.0	0	0	211536	13.00	NaN	S
887	female	19.0	0	0	112053	30.00	B42	S
888	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	male	26.0	0	0	111369	30.00	C148	C
890	male	32.0	0	0	370376	7.75	NaN	Q

```
[ ]: df.columns
```

```
[ ]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
          'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
          dtype='object')
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
```

```

9   Fare      891 non-null   float64
10  Cabin     204 non-null   object
11  Embarked  889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

```

```
[ ]: df.describe().T
```

```

[ ]:

```

	count	mean	std	min	25%	50%	75%	\
PassengerId	891.0	446.000000	257.353842	1.00	223.5000	446.0000	668.5	
Survived	891.0	0.383838	0.486592	0.00	0.0000	0.0000	1.0	
Pclass	891.0	2.308642	0.836071	1.00	2.0000	3.0000	3.0	
Age	714.0	29.699118	14.526497	0.42	20.1250	28.0000	38.0	
SibSp	891.0	0.523008	1.102743	0.00	0.0000	0.0000	1.0	
Parch	891.0	0.381594	0.806057	0.00	0.0000	0.0000	0.0	
Fare	891.0	32.204208	49.693429	0.00	7.9104	14.4542	31.0	


```

max
PassengerId  891.0000
Survived      1.0000
Pclass        3.0000
Age           80.0000
SibSp         8.0000
Parch         6.0000
Fare          512.3292

```

Data Cleaning

```
[ ]: #check for missing values
df.isnull().sum()
```

```

[ ]: PassengerId    0
Survived          0
Pclass            0
Name              0
Sex               0
Age              177
SibSp             0
Parch             0
Ticket           0
Fare              0
Cabin            687
Embarked          2
dtype: int64

```

```

[ ]: # Handle missing values (for simplicity, we'll drop the missing values in this
     ↪ example)
df1 =df.dropna(subset=['Age','Cabin'])

```

```
[ ]: # Confirm that there are no missing values anymore
df1.isnull().sum()
```

```
[ ]: PassengerId    0
      Survived      0
      Pclass       0
      Name         0
      Sex          0
      Age          0
      SibSp        0
      Parch        0
      Ticket       0
      Fare         0
      Cabin        0
      Embarked     2
      dtype: int64
```

```
[ ]: df1['Embarked'].fillna(df1['Embarked'].mode()[0], inplace=True)
```

<ipython-input-15-96f9b2e3057e>:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df1['Embarked'].fillna(df1['Embarked'].mode()[0], inplace=True)

```
[ ]: df1.shape
```

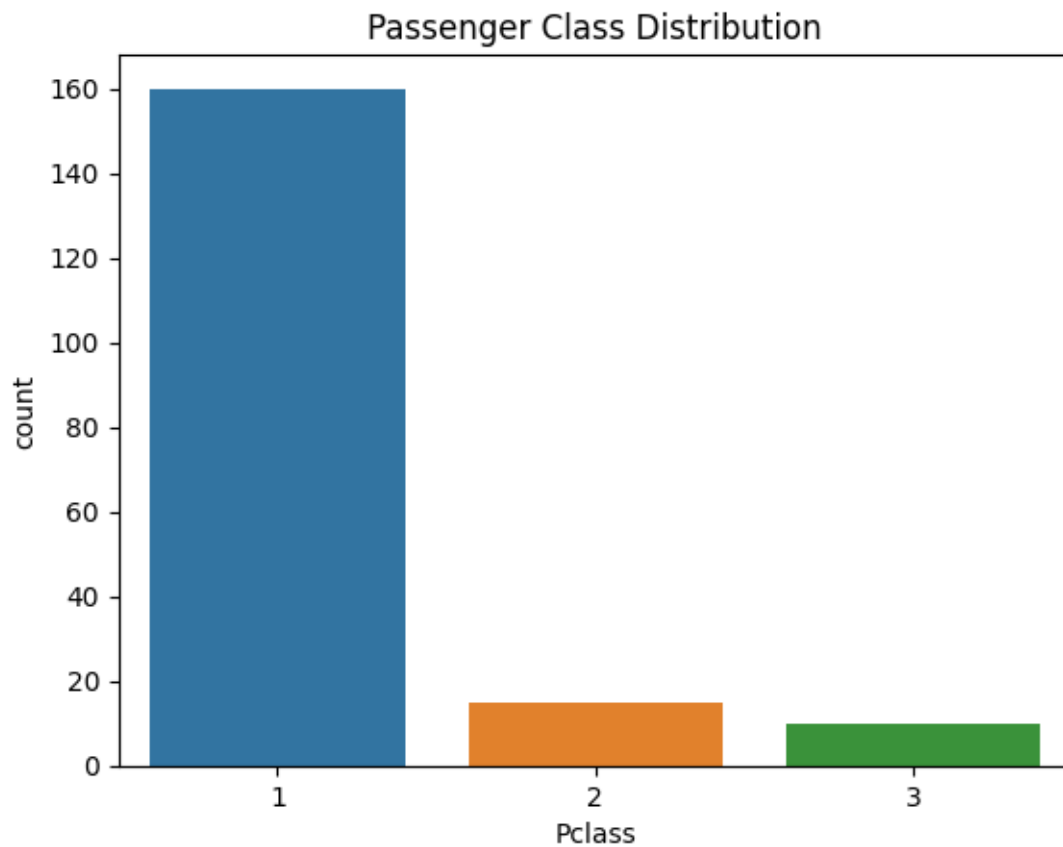
```
[ ]: (185, 12)
```

Visualization


```
[ ]: # Countplot for survival
      sns.countplot(x='Survived', data=df1)
      plt.title('Survival Count')
      plt.show()
```




```
[ ]: # Countplot for Pclass
sns.countplot(x='Pclass', data=df1)
plt.title('Passenger Class Distribution')
plt.show()
```




```
[ ]: # Boxplot for Age distribution by Pclass
sns.boxplot(x='Pclass', y='Age', data=df1)
plt.title('Age Distribution by Passenger Class')
plt.show()
```



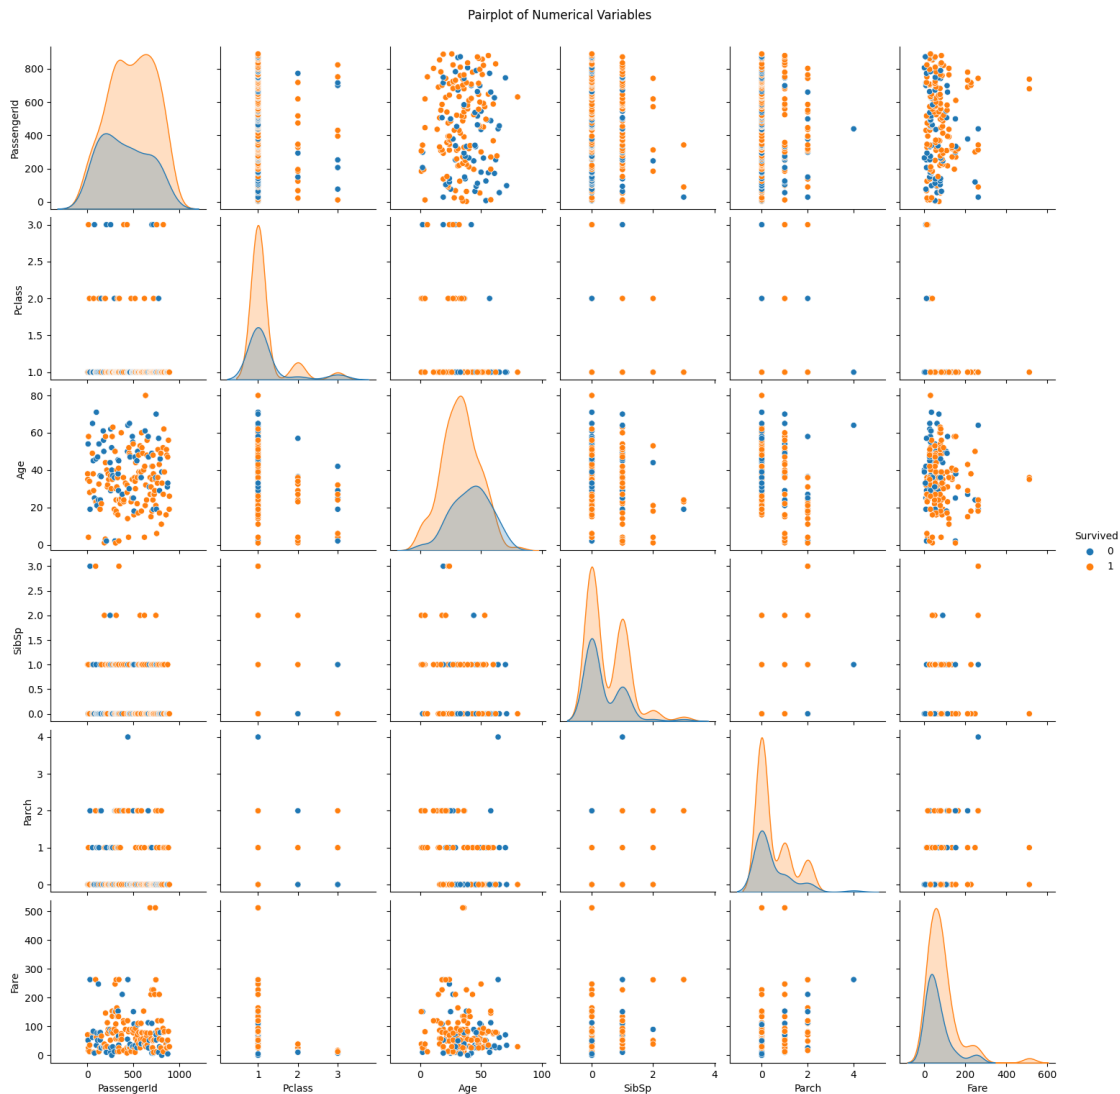
```
[ ]: # Countplot for Sex
sns.countplot(x='Sex', data=df1)
plt.title('Gender Distribution')
plt.show()
```

```
[ ]: # Countplot for Embarked
sns.countplot(x='Embarked', data=df1)
plt.title('Embarked Distribution')
plt.show()
```




```
[ ]: # Pairplot to explore relationships between numerical variables
sns.pairplot(df1, hue='Survived')
plt.suptitle('Pairplot of Numerical Variables', y=1.02)
plt.show()
```



```
[ ]: # Correlation heatmap
correlation_matrix = df1.corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap')
plt.show()
```


<ipython-input-23-fe033e8ba074>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_matrix = df1.corr()
```




Explore relationships and patterns


```
[ ]: # Survival by Passenger Class
sns.countplot(x='Pclass', hue='Survived', data=df1)
plt.title('Survival by Passenger Class')
plt.show()
```




```
[ ]: # Survival by Gender
sns.countplot(x='Sex', hue='Survived', data=df1)
plt.title('Survival by Gender')
plt.show()
```




```
[ ]: # Survival by Embarked Port
sns.countplot(x='Embarked', hue='Survived', data=df1)
plt.title('Survival by Embarked Port')
plt.show()
```



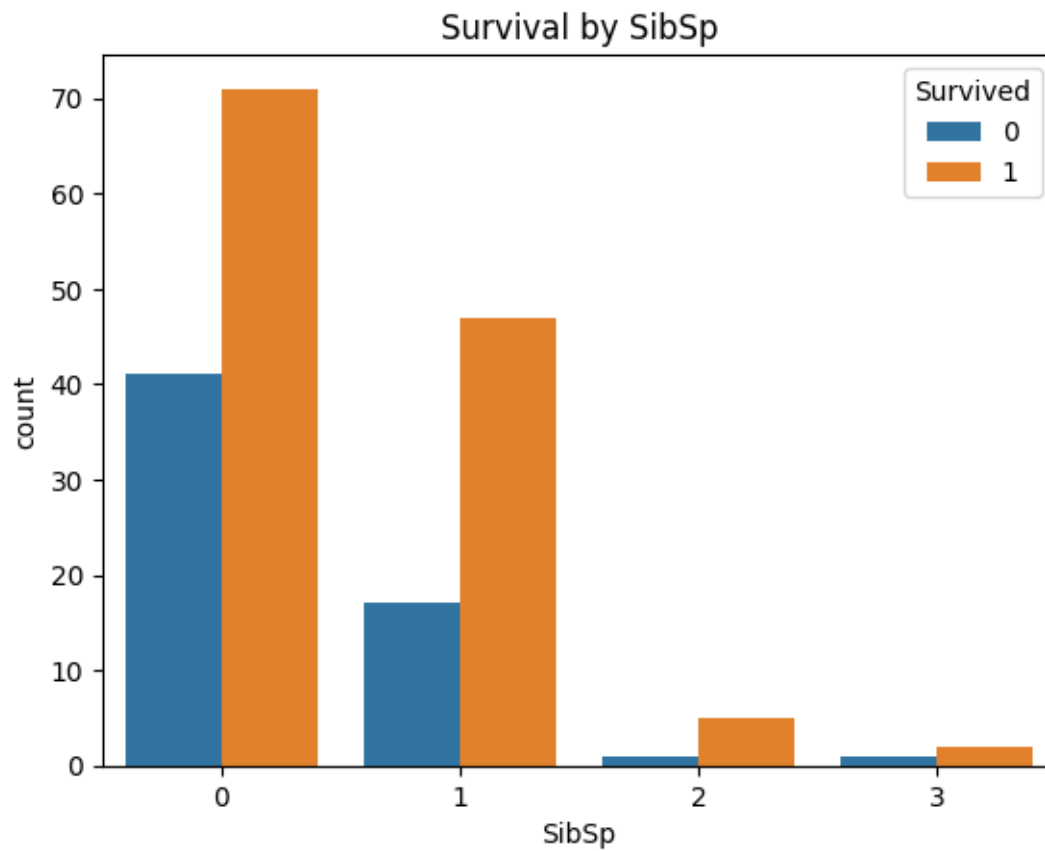
```
[ ]: # Age distribution by Survival
sns.histplot(x='Age', hue='Survived', data=df1, kde=True)
plt.title('Age Distribution by Survival')
plt.show()
```



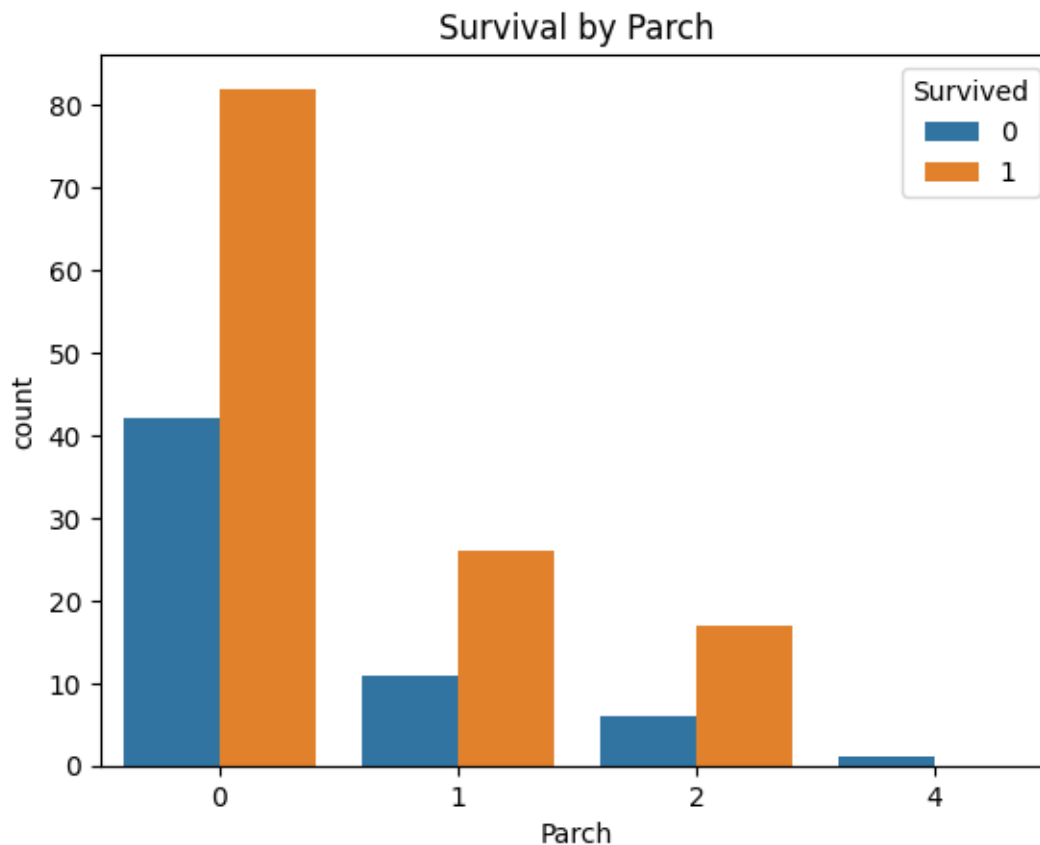
```
[ ]: # Fare distribution by Survival
sns.histplot(x='Fare', hue='Survived', data=df1, kde=True)
plt.title('Fare Distribution by Survival')
plt.show()
```

```
[ ]: # Survival by SibSp (Number of Siblings/Spouses Aboard)
sns.countplot(x='SibSp', hue='Survived', data=df1)
plt.title('Survival by SibSp')
plt.show()
```



```
[ ]: # Survival by Parch (Number of Parents/Children Aboard)
sns.countplot(x='Parch', hue='Survived', data=df1)
plt.title('Survival by Parch')
plt.show()
```



```
[ ]: df1.head()
```

```
[ ]:
  PassengerId  Survived  Pclass  \
1            2         1       1
3            4         1       1
6            7         0       1
10           11         1       3
11           12         1       1
```

```

                                Name      Sex  Age  SibSp  \
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0    1
3      Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0    1
6                        McCarthy, Mr. Timothy J    male  54.0    0
10                     Sandstrom, Miss. Marguerite Rut  female   4.0    1
11                     Bonnell, Miss. Elizabeth    female  58.0    0
```

```

   Parch  Ticket   Fare Cabin Embarked
1      0  PC 17599  71.2833   C85        C
3      0  113803  53.1000  C123        S
6      0   17463  51.8625  E46        S
```

10	1	PP 9549	16.7000	G6	S
11	0	113783	26.5500	C103	S

Label Encoding

```
[ ]: from sklearn.preprocessing import LabelEncoder
lb = LabelEncoder()
df1['Sex']=lb.fit_transform(df1['Sex'])
df1['Embarked']=lb.fit_transform(df1['Embarked'])
```

<ipython-input-32-591ec63db8c1>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1['Sex']=lb.fit_transform(df1['Sex'])
<ipython-input-32-591ec63db8c1>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1['Embarked']=lb.fit_transform(df1['Embarked'])
```

```
[ ]: #Train Test Split
# Select features and target variable
X = df1[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']]
y = df1['Survived']
```

```
[ ]: # Split the data into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪random_state=42)
```

Standardization

```
[ ]: from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
X_train = ss.fit_transform(X_train)
X_test = ss.transform(X_test)
```

Model building - Logistic Regression

```
[ ]: # Create and train a logistic regression model
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(random_state=33)
model.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(random_state=33)
```

```
[ ]: # Make predictions on the test set
y_pred = model.predict(X_test)
```

Evaluation Metrics

```
[ ]: # Evaluate the model
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
accuracy = accuracy_score(y_test, y_pred)*100
conf_matrix = confusion_matrix(y_test, y_pred)
classification_rep = classification_report(y_test, y_pred)

print(f'Accuracy: {accuracy:.4f}')
print(f'Confusion Matrix:\n{conf_matrix}')
print(f'Classification Report:\n{classification_rep}')
```

Accuracy: 75.6757

Confusion Matrix:

```
[[11  4]
 [ 5 17]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.69	0.73	0.71	15
1	0.81	0.77	0.79	22
accuracy			0.76	37
macro avg	0.75	0.75	0.75	37
weighted avg	0.76	0.76	0.76	37

In this example, we use logistic regression as a classifier. We preprocess the data by handling missing values, encoding categorical variables using Label Encoding, and standardizing numerical features. Finally, we train the model, make predictions on the test set, and evaluate its performance using accuracy, confusion matrix, and classification report.

```
[ ]: #BY HARI
```

```
[ ]: #Happy coding
```