

nz5qoe4ab

December 18, 2023

Task-04

- Analyze and visualize sentiment patterns in social media data to understand public opinion and attitudes towards specific topics or brands.

Sample Dataset :- <https://www.kaggle.com/datasets/jp797498e/twitter-entity-sentiment-analysis>

1 Description

About Dataset: this is the Twitter Sentiment Analysis Dataset.

Overview: This is an entity-level sentiment analysis dataset of twitter. Given a message and an entity, the task is to judge the sentiment of the message about the entity. There are three classes in this dataset: Positive, Negative and Neutral. We regard messages that are not relevant to the entity (i.e. Irrelevant) as Neutral.

Problem Statement: A Twitter sentiment analysis uses NLP and ML models to classify tweets into negative, positive or neutral emotions.

#Table of contents: * Import Modules * Exploratory data analysis(EDA) * Data cleaning * Data Visualization of Target Variables * Preprocessed text * Machine Learning Model

2 Import Modules

```
[ ]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import
    classification_report, accuracy_score, confusion_matrix
from mlxtend.plotting import plot_confusion_matrix
from sklearn.naive_bayes import MultinomialNB
import nltk
from nltk.sentiment.vader import SentimentIntensityAnalyzer
```

```
[ ]: # Load the dataset
df=pd.read_csv('twitter_training.csv')
```

```
[ ]: df
```

```
[ ]:
      2401  Borderlands  Positive  \
0      2401  Borderlands  Positive
1      2401  Borderlands  Positive
2      2401  Borderlands  Positive
3      2401  Borderlands  Positive
4      2401  Borderlands  Positive
...      ...      ...      ...
74676  9200      Nvidia  Positive
74677  9200      Nvidia  Positive
74678  9200      Nvidia  Positive
74679  9200      Nvidia  Positive
74680  9200      Nvidia  Positive

      im getting on borderlands and i will murder you all ,
0      I am coming to the borders and I will kill you...
1      im getting on borderlands and i will kill you ...
2      im coming on borderlands and i will murder you...
3      im getting on borderlands 2 and i will murder ...
4      im getting into borderlands and i can murder y...
...
74676  Just realized that the Windows partition of my...
74677  Just realized that my Mac window partition is ...
74678  Just realized the windows partition of my Mac ...
74679  Just realized between the windows partition of...
74680  Just like the windows partition of my Mac is l...

[74681 rows x 4 columns]
```

```
[ ]: # Exploratory data analysis(EDA)
```

```
[ ]: df.columns=['id','country','label','text']
```

```
[ ]: # shallow copy
df2=df.copy()
```

```
[ ]: #shape of a DataFrame.
df.shape
```

```
[ ]: (74681, 4)
```

```
[ ]: # displays the top rows of a DataFrame
df.head()
```

```
[ ]:      id      country      label \
0  2401  Borderlands  Positive
1  2401  Borderlands  Positive
2  2401  Borderlands  Positive
3  2401  Borderlands  Positive
4  2401  Borderlands  Positive

                                text
0  I am coming to the borders and I will kill you...
1  im getting on borderlands and i will kill you ...
2  im coming on borderlands and i will murder you...
3  im getting on borderlands 2 and i will murder ...
4  im getting into borderlands and i can murder y...
```

```
[ ]: #shows the bottom rows
df.tail()
```

```
[ ]:      id country      label \
74676  9200  Nvidia  Positive
74677  9200  Nvidia  Positive
74678  9200  Nvidia  Positive
74679  9200  Nvidia  Positive
74680  9200  Nvidia  Positive

                                text
74676  Just realized that the Windows partition of my...
74677  Just realized that my Mac window partition is ...
74678  Just realized the windows partition of my Mac ...
74679  Just realized between the windows partition of...
74680  Just like the windows partition of my Mac is l...
```

```
[ ]: # specific rows of a DataFrame ( "integer location" Method)
df.iloc[100:200]
```

```
[ ]:      id      country      label \
100  2417  Borderlands      Negative
101  2418  Borderlands  Irrelevant
102  2418  Borderlands  Irrelevant
103  2418  Borderlands  Irrelevant
104  2418  Borderlands  Irrelevant
..    ...      ...      ...
195  2433  Borderlands      Neutral
196  2433  Borderlands      Neutral
197  2434  Borderlands      Negative
198  2434  Borderlands      Negative
199  2434  Borderlands      Negative
```

```

                                text
100 Grounded almost was pretty cool even despite t...
101 Appreciate the (sonic) concepts / praxis Valen...
102 Appreciate the (sound) concepts / practices th...
103 Evaluate the (sound) concepts / concepts of Va...
104 Appreciate the (sonic) concepts / praxis Valen...
..
195 i then enter in that gunner seat and i fear fo...
196     i enter that gunner seat and i fear for a life
197         fuck it . pic.twitter.com/Wav1bacr5j
198         Fuck it. pic.twitter.com / Wav1bacr5j
199         fuck it. pic.wikipedia.org / Wav1bacr5j

[100 rows x 4 columns]

```

```
[ ]: # prints information about the DataFrame.
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 74681 entries, 0 to 74680
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   id           74681 non-null  int64
1   country      74681 non-null  object
2   label        74681 non-null  object
3   text         73995 non-null  object
dtypes: int64(1), object(3)
memory usage: 2.3+ MB

```

```
[ ]: # Display (string) columns in the summary statistics.
df.describe(include=object)
```

```

[ ]:
count          country      label      text
unique              32         4  69490
top    TomClancysRainbowSix  Negative
freq              2400     22542     172

```

3 Data cleaning

```
[ ]: # To check for duplicate values in a DataFrame
df.duplicated().sum()
```

```
[ ]: 2700
```

```
[ ]: # Remove duplicates based on all columns
df.drop_duplicates(inplace=True)
```

```
[ ]: # again To check for duplicate values in a DataFrame again
df.duplicated().sum()
```

```
[ ]: 0
```

```
[ ]: # The number of missing values in the dataset.
df.isnull().sum()
```

```
[ ]: id          0
country        0
label          0
text          326
dtype: int64
```

```
[ ]: # Drop rows with NaN values in-place
df.dropna(inplace=True)
```

```
[ ]: # our dataset remove null values
df.isnull().any()
```

```
[ ]: id          False
country        False
label          False
text           False
dtype: bool
```


#Data Visualization of Target Variables

```
[ ]: # Check unique target values
df['label'].value_counts()
```

```
[ ]: Negative      21698
Positive        19712
Neutral         17708
Irrelevant      12537
Name: label, dtype: int64
```

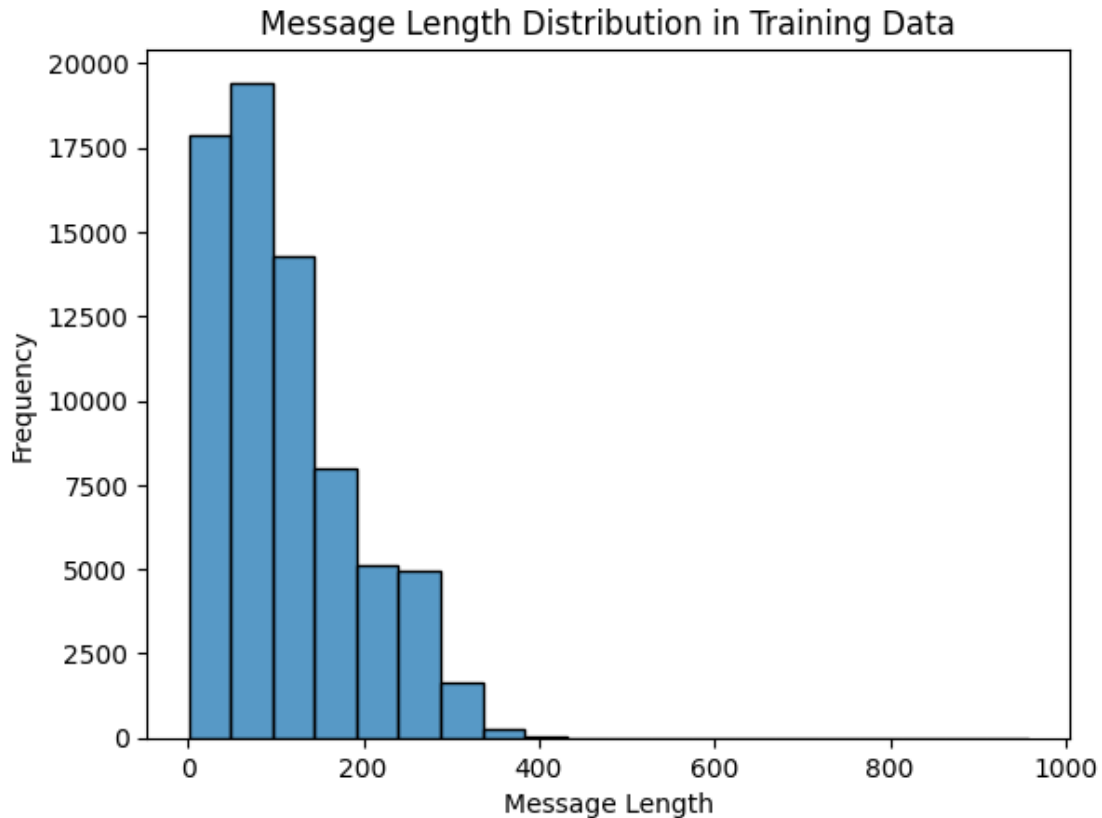
```
[ ]: sns.countplot(x=df['label'])
```

```
[ ]: <Axes: xlabel='label', ylabel='count'>
```



```
[ ]: # Calculate the length of each message
message_length=(df['text']).apply(len)
sns.histplot(x=message_length,bins=20)
plt.title('Message Length Distribution in Training Data')
plt.ylabel('Frequency')
plt.xlabel('Message Length')
```

```
[ ]: Text(0.5, 0, 'Message Length')
```



#Preprocessed text

```
[ ]: import spacy
# load english language model and create nlp object from it
nlp = spacy.load("en_core_web_sm")
# use this utility function to get the preprocessed text data
def preprocess(text):
    # remove stop words and lemmatize the text
    doc = nlp(text)
    filtered_tokens = []
    for token in doc:
        if token.is_stop or token.is_punct:
            continue
        filtered_tokens.append(token.lemma_)
    return " ".join(filtered_tokens)
```

```
[ ]: df['Preprocessed text'] = df['text'].apply(preprocess)
```

```
[ ]: lb=LabelEncoder()
df['label']=lb.fit_transform(df['label'])
```

```
[ ]: df
```

```
[ ]:      id      country  label  \
0      2401  Borderlands      3
1      2401  Borderlands      3
2      2401  Borderlands      3
3      2401  Borderlands      3
4      2401  Borderlands      3
...
74676  9200      Nvidia      3
74677  9200      Nvidia      3
74678  9200      Nvidia      3
74679  9200      Nvidia      3
74680  9200      Nvidia      3

                                text  \
0      I am coming to the borders and I will kill you...
1      im getting on borderlands and i will kill you ...
2      im coming on borderlands and i will murder you...
3      im getting on borderlands 2 and i will murder ...
4      im getting into borderlands and i can murder y...
...
74676  Just realized that the Windows partition of my...
74677  Just realized that my Mac window partition is ...
74678  Just realized the windows partition of my Mac ...
74679  Just realized between the windows partition of...
74680  Just like the windows partition of my Mac is l...

                                Preprocessed text
0                                come border kill
1                                m get borderland kill
2                                m come borderland murder
3                                m get borderland 2 murder
4                                m get borderland murder
...
74676  realize Windows partition Mac like 6 year Nvid...
74677  realize Mac window partition 6 year Nvidia dri...
74678  realize window partition Mac 6 year Nvidia dri...
74679  realize window partition Mac like 6 year Nvidi...
74680  like window partition Mac like 6 year driver i...

[71655 rows x 5 columns]
```

```
[ ]: tv=TfidfVectorizer()
df_tv=tv.fit_transform(df['Preprocessed text'])
```

```
[ ]: print(df_tv)
```


(0, 14186)	0.5019686782389964
(0, 4300)	0.7503332981844422
(0, 5882)	0.43014809973153667
(1, 4303)	0.6308352317883091
(1, 10718)	0.4731922339217186
(1, 14186)	0.6149276543551802
(2, 16730)	0.7359220742014858
(2, 4303)	0.519630312809822
(2, 5882)	0.4340541886817236
(3, 16730)	0.7497229075893237
(3, 4303)	0.5293750013057333
(3, 10718)	0.3970864765115596
(4, 16730)	0.7497229075893237
(4, 4303)	0.5293750013057333
(4, 10718)	0.3970864765115596
(5, 16356)	0.32986143201396134
(5, 5868)	0.0950308449908003
(5, 25306)	0.12371465037450177
(5, 18780)	0.12279967472353039
(5, 8680)	0.17199301599436456
(5, 6478)	0.31519414526267836
(5, 26163)	0.2882003846504435
(5, 12710)	0.23515040647542382
(5, 17993)	0.2103819690143733
(5, 18508)	0.17463994232150065
:	:
(71652, 17401)	0.335661757431383
(71652, 12602)	0.28873546946764583
(71652, 20209)	0.3195397101596675
(71652, 27556)	0.21093083092118967
(71653, 18390)	0.41917259340568874
(71653, 17512)	0.20004410985809554
(71653, 26966)	0.30984190903656667
(71653, 8064)	0.2857211695158495
(71653, 4956)	0.2711117868352008
(71653, 7524)	0.31837801158630585
(71653, 15399)	0.32982978949582387
(71653, 17401)	0.2933694892495072
(71653, 12602)	0.25235575793365683
(71653, 20209)	0.2792787664637086
(71653, 10264)	0.19437024500723696
(71653, 27556)	0.18435424579749274
(71653, 14875)	0.15320656386788417
(71654, 18390)	0.48735842343812535
(71654, 26966)	0.36024317113922943
(71654, 8064)	0.3321987670681811
(71654, 15399)	0.3834824335856304
(71654, 17401)	0.34109122116939317

```
(71654, 12602)      0.29340588165087583
(71654, 27556)      0.21434272182731726
(71654, 14875)      0.3562566379656403
```

```
[ ]: x_train, x_test, y_train, y_test = train_test_split(df_tv,
↳df['label'],test_size=0.2, random_state=42)
```

```
[ ]: x_test.shape
```

```
[ ]: (14331, 28054)
```

```
[ ]: y_test.shape
```

```
[ ]: (14331,)
```

```
#Machine Learning Model
```

```
##Naive bayes
```

```
[ ]: nb=MultinomialNB()
nb.fit(x_train,y_train)
y_pred_nb=nb.predict(x_test)
print('classification_report:\n',classification_report(y_test,y_pred_nb))
print('accuracy:',accuracy_score(y_test,y_pred_nb)*100)
print('Error value',np.mean(y_pred_nb!=y_test)*100)
print('confusion_matrix\n',confusion_matrix(y_test,y_pred_nb))
```

```
classification_report:
```

	precision	recall	f1-score	support
0	0.95	0.44	0.61	2455
1	0.66	0.89	0.76	4433
2	0.82	0.64	0.72	3532
3	0.69	0.79	0.74	3911
accuracy			0.72	14331
macro avg	0.78	0.69	0.71	14331
weighted avg	0.76	0.72	0.72	14331

```
accuracy: 72.40946200544275
```


```
Error value 27.590537994557252
```

```
confusion_matrix
```

```
[[1092 669 142 552]
 [ 17 3929 168 319]
 [ 23 754 2265 490]
 [ 17 615 188 3091]]
```

```
[ ]: plot_confusion_matrix(confusion_matrix(y_test,y_pred_nb))
```

```
[ ]: (<Figure size 640x480 with 1 Axes>,
      <Axes: xlabel='predicted label', ylabel='true label'>)
```



3.1 Logistic regression

```
[ ]: lg=LogisticRegression()
lg.fit(x_train,y_train)
y_pred_lg=lg.predict(x_test)
print('classification_report:\n',classification_report(y_test,y_pred_lg))
print('accuracy:',accuracy_score(y_test,y_pred_lg)*100)
print('Error value',np.mean(y_pred_lg!=y_test)*100)
print('confusion_matrix\n',confusion_matrix(y_test,y_pred_lg))
```

```
classification_report:
      precision    recall  f1-score   support

0         0.81      0.66      0.73      2455
1         0.80      0.83      0.81      4433
2         0.73      0.76      0.75      3532
3         0.75      0.79      0.77      3911
```

accuracy			0.77	14331
macro avg	0.77	0.76	0.76	14331
weighted avg	0.77	0.77	0.77	14331

accuracy: 77.13348684669597

Error value 22.866513153304027

confusion_matrix

```
[[1616 263 246 330]
 [ 119 3666 324 324]
 [ 127 352 2690 363]
 [ 131 293 405 3082]]
```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>


Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

n_iter_i = _check_optimize_result(

```
[ ]: plot_confusion_matrix(confusion_matrix(y_test,y_pred_lg))
```

```
[ ]: (<Figure size 640x480 with 1 Axes>,
      <Axes: xlabel='predicted label', ylabel='true label'>)
```



#VADER Sentiment Analysis

```
[ ]: nltk.download('vader_lexicon')
      sid = SentimentIntensityAnalyzer()
```

```
[nltk_data] Downloading package vader_lexicon to /root/nltk_data...
[nltk_data] Package vader_lexicon is already up-to-date!
```

```
[ ]: # Function to get sentiment scores for a given text
      def get_sentiment_scores(text):
          sentiment_scores = sid.polarity_scores(text)
          return sentiment_scores
```

```
[ ]: df3=df.copy()
```

```
[ ]: # Apply the sentiment analysis function to the 'text' column and create new
      ↪ columns for scores
      df3['sentiment_scores'] = df3['text'].apply(get_sentiment_scores)
```

```
[ ]: df3['sentiment_scores']
```

```
[ ]: 0      {'neg': 0.343, 'neu': 0.657, 'pos': 0.0, 'comp...
      1      {'neg': 0.37, 'neu': 0.63, 'pos': 0.0, 'compou...
      2      {'neg': 0.37, 'neu': 0.63, 'pos': 0.0, 'compou...
      3      {'neg': 0.343, 'neu': 0.657, 'pos': 0.0, 'comp...
      4      {'neg': 0.37, 'neu': 0.63, 'pos': 0.0, 'compou...

      ...
      74676   {'neg': 0.086, 'neu': 0.817, 'pos': 0.097, 'co...
      74677   {'neg': 0.104, 'neu': 0.896, 'pos': 0.0, 'comp...
      74678   {'neg': 0.091, 'neu': 0.909, 'pos': 0.0, 'comp...
      74679   {'neg': 0.074, 'neu': 0.842, 'pos': 0.084, 'co...
      74680   {'neg': 0.09, 'neu': 0.728, 'pos': 0.182, 'com...
      Name: sentiment_scores, Length: 71655, dtype: object
```

```
[ ]: # Extract individual sentiment scores into separate columns
df3['compound'] = df3['sentiment_scores'].apply(lambda x: x['compound'])
df3['positive'] = df3['sentiment_scores'].apply(lambda x: x['pos'])
df3['neutral'] = df3['sentiment_scores'].apply(lambda x: x['neu'])
df3['negative'] = df3['sentiment_scores'].apply(lambda x: x['neg'])
```

```
[ ]: df3['sentiment'] = df3['compound'].apply(lambda x: 'Positive' if x >= 0.05 else
      ↪('Negative' if x <= -0.05 else 'Neutral'))
```

```
[ ]: df3
```

```
[ ]:      id      country  label  \
0      2401  Borderlands      3
1      2401  Borderlands      3
2      2401  Borderlands      3
3      2401  Borderlands      3
4      2401  Borderlands      3
...      ...
74676   9200      Nvidia      3
74677   9200      Nvidia      3
74678   9200      Nvidia      3
74679   9200      Nvidia      3
74680   9200      Nvidia      3

      text  \
0      I am coming to the borders and I will kill you...
1      im getting on borderlands and i will kill you ...
2      im coming on borderlands and i will murder you...
3      im getting on borderlands 2 and i will murder ...
4      im getting into borderlands and i can murder y...
...      ...
74676   Just realized that the Windows partition of my...
74677   Just realized that my Mac window partition is ...
74678   Just realized the windows partition of my Mac ...
```

```

74679 Just realized between the windows partition of...
74680 Just like the windows partition of my Mac is l...

```

```

                                Preprocessed text \
0                                come border kill
1                                m get borderland kill
2                                m come borderland murder
3                                m get borderland 2 murder
4                                m get borderland murder
...
74676 realize Windows partition Mac like 6 year Nvid...
74677 realize Mac window partition 6 year Nvidia dri...
74678 realize window partition Mac 6 year Nvidia dri...
74679 realize window partition Mac like 6 year Nvidi...
74680 like window partition Mac like 6 year driver i...

```

```

                                sentiment_scores compound positive \
0      {'neg': 0.343, 'neu': 0.657, 'pos': 0.0, 'comp... -0.6908      0.000
1      {'neg': 0.37, 'neu': 0.63, 'pos': 0.0, 'compou... -0.6908      0.000
2      {'neg': 0.37, 'neu': 0.63, 'pos': 0.0, 'compou... -0.6908      0.000
3      {'neg': 0.343, 'neu': 0.657, 'pos': 0.0, 'comp... -0.6908      0.000
4      {'neg': 0.37, 'neu': 0.63, 'pos': 0.0, 'compou... -0.6908      0.000
...
74676 {'neg': 0.086, 'neu': 0.817, 'pos': 0.097, 'co...      0.0772      0.097
74677 {'neg': 0.104, 'neu': 0.896, 'pos': 0.0, 'comp... -0.2960      0.000
74678 {'neg': 0.091, 'neu': 0.909, 'pos': 0.0, 'comp... -0.2960      0.000
74679 {'neg': 0.074, 'neu': 0.842, 'pos': 0.084, 'co...      0.0772      0.084
74680 {'neg': 0.09, 'neu': 0.728, 'pos': 0.182, 'com...      0.3687      0.182

```

```

                                neutral negative sentiment
0                                0.657      0.343 Negative
1                                0.630      0.370 Negative
2                                0.630      0.370 Negative
3                                0.657      0.343 Negative
4                                0.630      0.370 Negative
...
74676      ...      ...      Positive
74677      0.896      0.104 Negative
74678      0.909      0.091 Negative
74679      0.842      0.074 Positive
74680      0.728      0.090 Positive

```

```
[71655 rows x 11 columns]
```

```
[ ]: df3['sentiment']
```

```
[ ]: 0      Negative
      1      Negative
      2      Negative
      3      Negative
      4      Negative
      ...
      74676   Positive
      74677   Negative
      74678   Negative
      74679   Positive
      74680   Positive
      Name: sentiment, Length: 71655, dtype: object
```

```
[ ]: # by HARI
```

```
[ ]: #HAPPY CODING!!!
```